Vinay Kudari

seeking a team that thrives on complex problems and candid, idea-driven dialogue

vinay.kudari30@gmail.com
+1-469-943-6778

Relevant skills

Python, SQL, *FastAPI, LangGraph, PySpark, PyTorch, Hugging Face, Kubernetes, Kafka, Terraform*, **GCP, AWS**

Experience [6+ yrs]

Apple — Cupertino, CA
Machine Learning Engineer | ***automation, mcp, tool-calling, agents*** — *July 2024 - Present*
- Architected an MCP (model context protocol) enabled asynchronous task processor that securely executes agent driven jobs in isolated, ephemeral Kubernetes pods at scale
- Implemented a multi-agent natural language query planner (intent parsing → schema mapping → rewrite/optimize) generating low-latency, resource-efficient SQL
- Automated end-to-end document review and database validation via a multi-agent system, replacing manual "read–query–verify–sign" workflows and cutting review labor substantially
- Engineering an agentic system that correlates transactional flows across SAP modules, detects volumetric outliers, enriches root-cause context, and files tickets

Gap Inc — San Francisco, CA
Software Engineer | ***java, spring, graphql, azure, kafka, kubernetes, llm*** — *July 2023 - June 2024*
- Spearheaded a real-time personalization engine development, leveraging extensive datasets to optimize user experiences and drive significant engagement and sales enhancements.
- Engineering a low-latency LLM-based information retrieval system to extract data from multi-source document corpus

Amazon Web Services — Palo Alto, CA
Software Development Engineer | ***spark, athena, dynamo, glue, lambda, cdk*** — *Dec 2022 - June 2023*
- Designed and published an internal query builder python package (ZenQL) tailored for AWS Healthlake
- Designed an event-driven server-less system using AWS services, allowing for parallel execution of tasks, resulting in seamless scaling to handle concurrent requests and improved processing throughput
- Designed and optimized complex SQL queries in Athena resulting in a 40% reduction in query execution times
- Architected near real time data processing system using AWS CDK, leveraging PySpark for ETL transforms on multi-terabyte healthcare data sets

Playstation — San Francisco, CA
Software Development Intern | ***kafka, concurrency, async, no-sql, cassandra*** — *May 2022 - Dec 2022*
- Worked on auto follow feature for PS5 Explore Hub. Message queues and REST endpoints were developed
- Learned custom framework and leveraged asynchronous programming pattern to optimize the latency by 60%
- Developed a cross platform CLI tool to automate generation of grafana analytical dashboards from logs

Brave Orbit (early stage startup) — Capetown, South Africa
Founding engineer | ***rest-api, django, flutter, pytest, redis, docker, gcp*** — *Jan 2020 - Aug 2021*
- Programmed a cross-platform app using BLoC pattern, implemented database, network and state management layers
- Engineered HIPAA compliant scalable async push notification service for medication reminder application
- Developed a task scheduler for clinical surveys that delivers email, SMS alerts and gathers response on a recurring basis
- Architected a custom ERP system with multi-product substitutions and combo-product life cycle tracking features

Accenture — Bangalore, India
Big Data Developer | ***redshift, pyspark, emr, etl, s3*** — *Sep 2018 - Dec 2019*
- Performed time series analysis on drug assay data to identify raw materials that effect the drug quality
- Designed ETL data validation pipeline using micro service architecture; Manual testing resources were cut by 70%
- Developed PL/SQL procedures to generate parent child hierarchies using temporal drug life cycle data

Projects

- **Friend (say hi @mydawg_bot on telegram)**: Building a multi-agent AI companion with human-like memory simulation, capable of personalized conversations, context-aware recall, and tool use
- **Stocklerts**: Built a self-evolving system that analyzes real-time news to recommend stocks for day trading; adapts prompts weekly based on ground-truth top performers to refine future predictions
- **Open-source contributions**: Contributor to HuggingFace Datasets official repository
- **Distributed key-value store**: Developed redis-like distributed key-value store with a focus on achieving strong consistency and fault tolerance using the RAFT consensus algorithm
- **Unordered image stitching and deghosting**: Estimated approximate overlap area by minimum variance technique on binarized images; stitched images based on binary tree model; removed deghosting from the overlap areas by selecting regions based on minimum local gradients

Education

University at Buffalo [3.9/4.0 GPA] — Buffalo, NY
Master of Science in Computer Science (AI/ML Specialization) — *Sep 2021 - Dec 2022*
Research: *Finding inconsistency between text and related images, advised by Prof. David Doermann*

National Institute of Technology, Durgapur [3.3/4.0 GPA] — West Bengal, India
Bachelor of Technology in Computer Science and Engineering — *July 2014 - May 2018*